

# Systematic Dynamic Race Condition Detection

*more clever than “slaughter cho2” since 2011.*

Ben Blum (bblum@andrew.cmu.edu)

Carnegie Mellon University - 15-410

2013, April 14

# Outline

## **Theory: Seeing race conditions in a new way**

- ▶ Case study (example)
- ▶ The execution tree
- ▶ Decision points

## **Technique: Systematic testing**

- ▶ Requirements
- ▶ Challenges and feasibility

## **Tool: Landslide**

- ▶ How it works
- ▶ The user-tool relationship
- ▶ User study (that's you!)

# Case Study

```
int thread_fork()
{
    thread_t *child = construct_new_thread();
    add_to_runqueue(child);
    return child->tid;
}
```

# Decision Points (“good” case)

| Thread 1                          | Thread 2            |                        |
|-----------------------------------|---------------------|------------------------|
| <code>spawn_new_thread</code>     |                     |                        |
| <code>add_to_runqueue</code>      |                     | (new thread)           |
| <code>return child-&gt;tid</code> |                     |                        |
|                                   | <code>vanish</code> |                        |
|                                   | (TCB gets freed)    | (voluntary reschedule) |

# Decision Points (race condition)

| Thread 1          | Thread 2         |
|-------------------|------------------|
| spawn_new_thread  |                  |
| add_to_runqueue   |                  |
|                   | vanish           |
|                   | (TCB gets freed) |
| return child->tid |                  |

(new thread + preempted)

(voluntary reschedule)

(bad!)

# Testing Mechanisms

**Stress testing:** slaughter cho2 and friends


- ▶ Attempting to exercise as many interleavings as practical
- ▶ Exposes race conditions at random
  - ▶ “If a preemption occurs at just the right time...”
- ▶ Cryptic panic messages or machine reboots

What if...

- ▶ Make educated guesses about when to preempt
- ▶ Preempt enough times to run *every single* interleaving
- ▶ Tell the story of what *actually happened*.
- ▶ Overlook fewer bugs!

A different way of looking at race conditions. . .

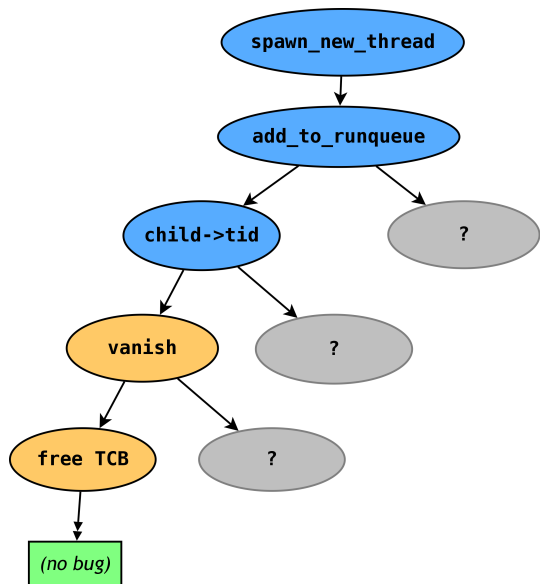
# Execution Tree



| Thread 1                          | Thread 2                      |
|-----------------------------------|-------------------------------|
| <code>spawn_new_thread</code>     |                               |
| <code>add_to_runqueue</code>      |                               |
| <code>return child-&gt;tid</code> |                               |
|                                   | <code>vanish</code>           |
|                                   | <code>(TCB gets freed)</code> |

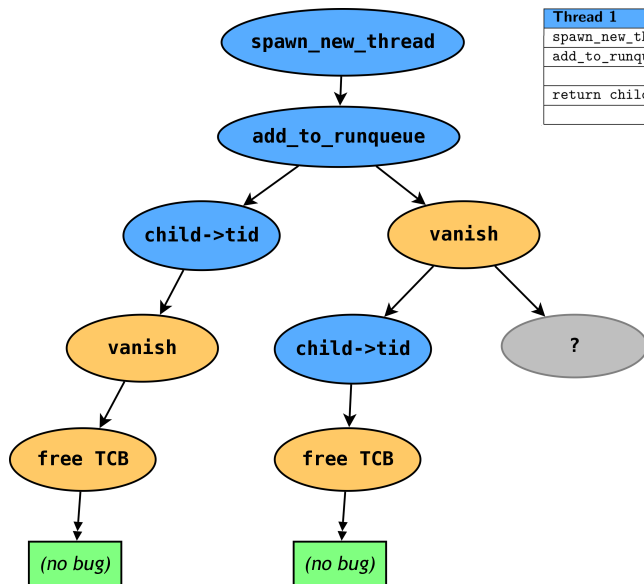


# Execution Tree



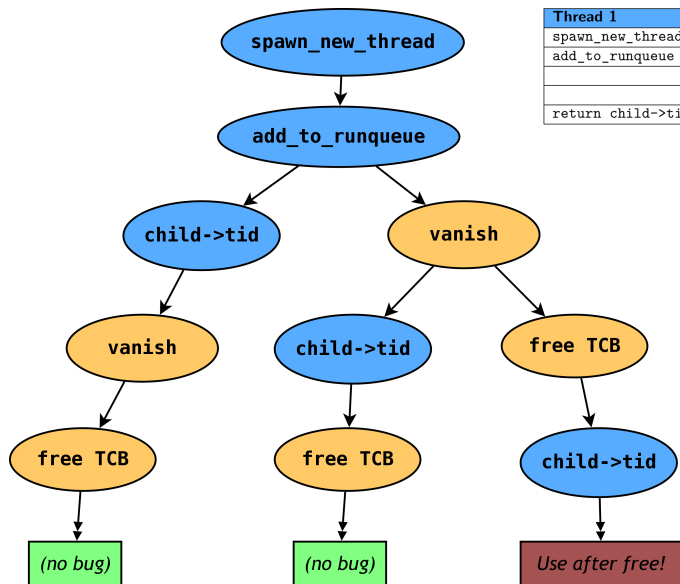
| Thread 1          | Thread 2         |
|-------------------|------------------|
| spawn_new_thread  |                  |
| add_to_runqueue   |                  |
| return child->tid |                  |
|                   | vanish           |
|                   | (TCB gets freed) |

# Execution Tree



| Thread 1          | Thread 2         |
|-------------------|------------------|
| spawn_new_thread  |                  |
| add_to_runqueue   |                  |
|                   | vanish           |
| return child->tid |                  |
|                   | (TCB gets freed) |

# Execution Tree



| Thread 1          | Thread 2                   |
|-------------------|----------------------------|
| spawn_new_thread  |                            |
| add_to_runqueue   |                            |
|                   | vanish<br>(TCB gets freed) |
| return child->tid |                            |



# Decision Points

A **decision point** is...

A code location where being preempted causes different behaviour.

- ▶ Intuitively: Somewhere that interesting interleavings can happen around.

Examples:

- ▶ A new thread becomes runnable
- ▶ Voluntary reschedule (e.g. `yield`, `cond_wait`)
- ▶ Synchronization primitives

# Systematic Testing

# Systematic Testing

**Systematic testing** is:

- ▶ Systematically enumerating different interleavings
  - ▶ Intuitively: Generate many “tabular execution traces”
- ▶ Exploring all branches in these trees
  - ▶ (by controlling scheduling decisions at decision points)
- ▶ In practice: Depth-first search of branches

# Execution Tree Exploration

Important point: When does a branch end?

- ▶ All threads run to completion, or
- ▶ A bug is detected

## **Backtracking:**

- ▶ Identify a decision point to choose differently
- ▶ Reset machine state and start over
- ▶ Replay test from the beginning, with different “decisions”

# More on Decision Points

Important point: What does “all possible interleavings” mean?

One extreme: Decide at every instruction

- ▶ *Good news*: Will find every possible race condition.
- ▶ *Bad news*: Runtime of test will be impossibly large.

Other extreme: Nothing is a decision point

- ▶ *Good news*: Test will finish quickly.
- ▶ *Bad news*: Only one execution was checked for bugginess.
- ▶ *Bad news*: No alternative interleavings explored.
  - ▶ Makes “no race found” a weak claim.



# More on Decision Points

Sweet spot: Insert a thread switch everywhere it “might matter”.

When do we fear being preempted?

- ▶ New threads becoming runnable (`fork`, `cond_signal`, etc)
  - ▶ Preemptions may cause it to run before we're ready
- ▶ Synchronization primitives (`mutex_lock`, etc)
  - ▶ If used improperly...
- ▶ Unprotected shared memory accesses
  - ▶ May result in data structure corruption

Finding the sweet spot is a joint effort between programmer and tool.  
(More on this later.)

# Controlling Scheduling Decisions

Control over sources of nondeterminism

- ▶ Device interrupts/input
  - ▶ Disk drivers: when disk reads finish
  - ▶ Ethernet drivers: when packets arrive
- ▶ To control thread switches in a 410 kernel, vary when clock ticks happen.

# Memory Interposition

In order to find use-after-free, need to know:

- ▶ When objects are `free()`d
- ▶ When threads access shared memory in the heap

Solution: Keep track of all memory events

- ▶ All calls to `malloc/free`
- ▶ All shared memory reads/writes

# State Space Explosion

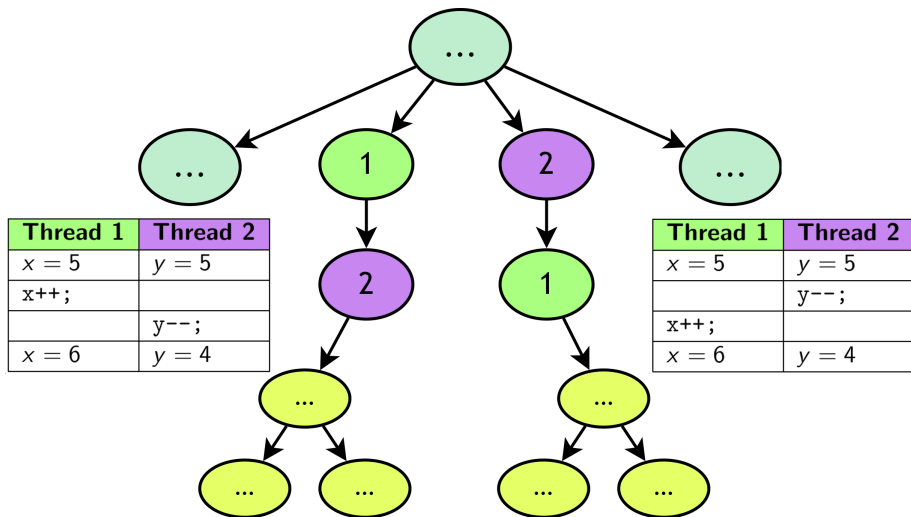
State spaces grow exponentially

- ▶ With  $d$  decision points,  $k$  runnable threads, size  $d^k$ .
- ▶ Threatens our ability to explore everything.
- ▶ Fortunately, some sequences result in identical states.

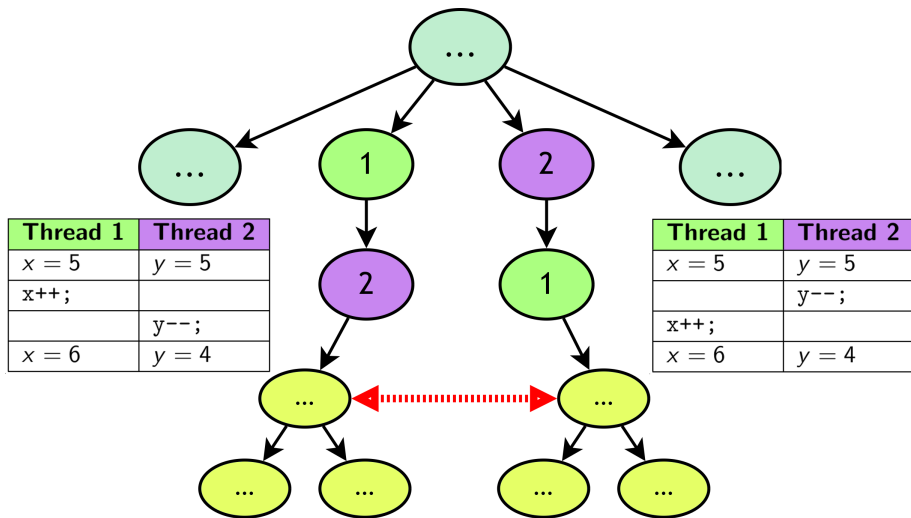
**Partial Order Reduction** can help.

- ▶ Complicated algorithm; ask me later for details.
- ▶ Intuitive explanation follows.

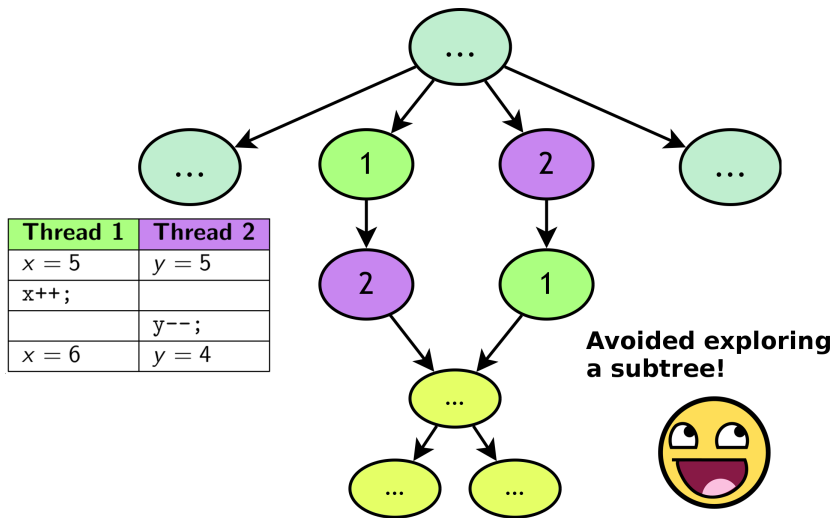
# State Space Explosion



# State Space Explosion



# State Space Explosion



# Landslide



# About The Project

5th year MS since June 2011 My MS thesis project (June 2011-2012)

Now I'm a Ph.D. student. . .

Working with Garth Gibson, Jiri Simsa

**Landslide:** Shows that your Pebbles are not as stable as you thought.

# Landslide in Simics

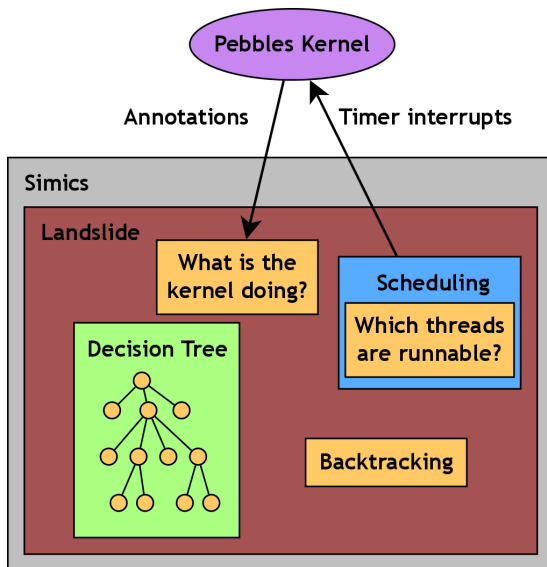
As a Simics module, Landslide knows:

- ▶ Every instruction the kernel executes
- ▶ Every memory address the kernel reads/writes

Artificially causes timer interrupts

Checkpointing/backtracking via Simics bookmarks (reverse execution)

# Anatomy



# Identifying Bugs

Landslide can *definitely discover*:

- ▶ Kernel panics
- ▶ Deadlock
- ▶ Use-after-free / double-free

Landslide can *reasonably suspect*:

- ▶ Memory leak
- ▶ Infinite loop (halting problem)

# Using Landslide

# In Which Ben Offers Help

This is something you can try!

Mutual benefit

- ▶ Landslide may help you find bugs
- ▶ You help Ben evaluate his research

# Keeping It Real

Finding race conditions is hard for humans.

It is hard for computer programs too.

Landslide is not an oracle.

# Annotating Your Kernel

## Step 1

```
int thread_fork()
{
    thread_t *child = spawn_new_thread();
    tell_landslide_forking();
    add_to_runqueue(child);
    tell_landslide_decide(); /* Interrupt me here! */
    return child->tid;
}
```

Your kernel needs to say when certain events happen:

- ▶ When do threads become runnable / descheduled?
- ▶ When does the scheduler switch threads?



# Configuring Landslide

## Step 2

```
# What test case to run?  
TEST_CASE="double_thread_fork"  
  
# The names of some important functions  
TIMER_WRAPPER="timer_handler_wrapper"  
CONTEXT_SWITCH="sched_switch"  
  
# What functions to pay attention to?  
within_function "thread_fork"  
within_function "vanish"
```

Edit `config.landslide` with some details and tweaks

Fill in two implementation-dependent C functions in Landslide ( $\leq 10$  lines)

# Configuring Decision Points

Landslide automatically identifies a minimal set of decision points.

- ▶ It might find bugs.
- ▶ It might overlook more fine-grained interleavings.

With help from you, it could find more.

- ▶ Optional annotation: `tell_landslide_decide()`
- ▶ Hints to where a context switch should be forced.
- ▶ Inside every call to `mutex_lock...`

# And Hopefully. . .

```

[SCHEDULE] thread 4 vanished
[SCHEDULE] switched threads 4 -> 3
[MEMORY] USE AFTER FREE - read from 0x0015a8f0 at eip 0x00104209
[MEMORY] Heap contents: {...}
[MEMORY] [0x15a8f0 | 4136] was allocated by TID3 at (...)
[MEMORY] and freed by TID4 at (...)
[BUG!] ***** A bug was found! *****
[BUG!] ***** Decision trace follows. *****
[BUG!] 1: 1347079 instructions, old 3 new 4, current 4
[BUG!] TID3 at 0x00105a10 in context_switch,
[BUG!] 0x001041f4 in thread_fork,
[BUG!] 0x0010362b in thread_fork_wrapper
[BUG!] 2: 1350725 instructions, old 4 new 3, current 3
[BUG!] TID3 at 0x00105a10 in context_switch,
[BUG!] 0x00104681 in yield,
[BUG!] 0x00104570 in vanish,
[BUG!] 0x00103708 in vanish_wrapper
[BUG!] Stack: TID3 at 0x00104209 in thread_fork,
[BUG!] 0x0010362b in thread_fork_wrapper
[BUG!] Total decision points 24, total backtracks 5
[BUG!] Average instrs/decision 16155, average branch depth 5

```

## Quick Demo

# Last Year's Results

- ▶ Worked with 5 groups
  - ▶ ...of which 4 put in enough time to make Landslide work
- ▶ **Investment:** about 2 to 3 hours
  - ▶ Average 112 minutes doing required instrumentation
  - ▶ Average 35 minutes configuring extra decision points & interpreting “found a bug” output

# Last Year's Results

- ▶ Worked with 5 groups
  - ▶ ...of which 4 put in enough time to make Landslide work
- ▶ **Investment:** about 2 to 3 hours
  - ▶ Average 112 minutes doing required instrumentation
  - ▶ Average 35 minutes configuring extra decision points & interpreting “found a bug” output
- ▶ **Return:** All 4 groups found bugs in their kernels.
  - ▶ All groups found deterministic bugs (e.g., use-after-free)
  - ▶ 2 groups found race conditions (i.e., Landslide needed to backtrack)

# In Which Ben Offers Help - Warning

**If you are already struggling, this will not “save” you.**

- ▶ False-negatives: not guaranteed to find races at all
- ▶ Research-quality: possibly difficult to integrate with your kernel
- ▶ Finishing the kernel project is more important.

# In Which Ben Offers Help - Warning

**If you are already struggling, this will not “save” you.**

- ▶ False-negatives: not guaranteed to find races at all
- ▶ Research-quality: possibly difficult to integrate with your kernel
- ▶ Finishing the kernel project is more important.

This is for you if:

- ▶ You have already submitted
- ▶ You are using late days, but willing to work over Carnival
- ▶ You are taking p3extra, but you already pass the hurdle
- ▶ You are looking for...
  - ▶ That “one pesky race”
  - ▶ A race that stress-testing missed
  - ▶ Or just familiarity with a new technique



# In Which Ben Offers Help

## Your kernel

- ▶ Must load the shell and run programs
- ▶ `fork`, `exec`, `vanish`, `wait`, `readline`
- ▶ Must never spin-wait (see hurdle form!)
- ▶ Should `assert()` important invariants
  - ▶ Think of `panic()` as `tell_landslide_bug()`

# In Which Ben Offers Help

User study this week, by appointment ([bblum@cs.cmu.edu](mailto:bblum@cs.cmu.edu))

Expect to spend:

- ▶ Up to 3 hours, just to try it out.
- ▶ 4-6 hours, if you find a bug and track it down.
- ▶ More, for multiple bugs or the truly dedicated...

Give feedback (intuitive? frustrating? found bugs?)

**Watch 410.announce for details!**

Questions?