

---

# Concurrency Testing with Data-Race Preemption Points

**Ben Blum**

**Garth Gibson**

PDL Retreat 2015

Parallel Data Laboratory

Carnegie Mellon University

# Outline

---

## **What is Systematic Testing?**

- Motivation
- Role of preemption points
- Challenges

## ***Iterative Deepening***

- Data race detection
- State space estimation
- Implementation in Landslide & Quicksand

## **Evaluation**

- Finding bugs in CMU OS student projects

## **Conclusion**

---

# Background: What is Systematic Testing?


# Example

---

```
int thread_fork()
{
    thread_t *child = spawn_new_thread();
    add_to_runqueue(child);
    return child->tid;
}
```

# Example

---

```
int thread_fork()  
{  
    thread_t *child = spawn_new_thread();  
    add_to_runqueue(child);  
    return child->tid;  “child” gets freed!  
}
```

- On exit, child's state is freed
- Forking thread does use-after-free
- Might return garbage instead of thread ID

# Example

---

Thread 1	Thread 2	
spawn_new_thread		
add_to_runqueue		(new thread)
return child->tid		(yield to child)
	exit	
	(TCB gets freed)	

Thread 1	Thread 2	
spawn_new_thread		
add_to_runqueue		(new thread + preempted)
	exit	
	(TCB gets freed)	(yield to parent)
return child->tid		(bad!)

# Systematic Testing

---

## Stress testing

- Common approach to finding concurrency bugs
- Relies on randomly-occurring context switches
- User-friendly but unpredictable/unbounded

# Systematic Testing

---

## Stress testing

- Common approach to finding concurrency bugs
- Relies on randomly-occurring context switches
- User-friendly but unpredictable/unbounded

## Systematic testing [*Godefroid '97*]

- Thread scheduling controlled by test framework
- Each iteration, test a different thread interleaving
- Goal: Exhaustively search “all possible” interleavings



# Systematic Testing

---

## Stress testing

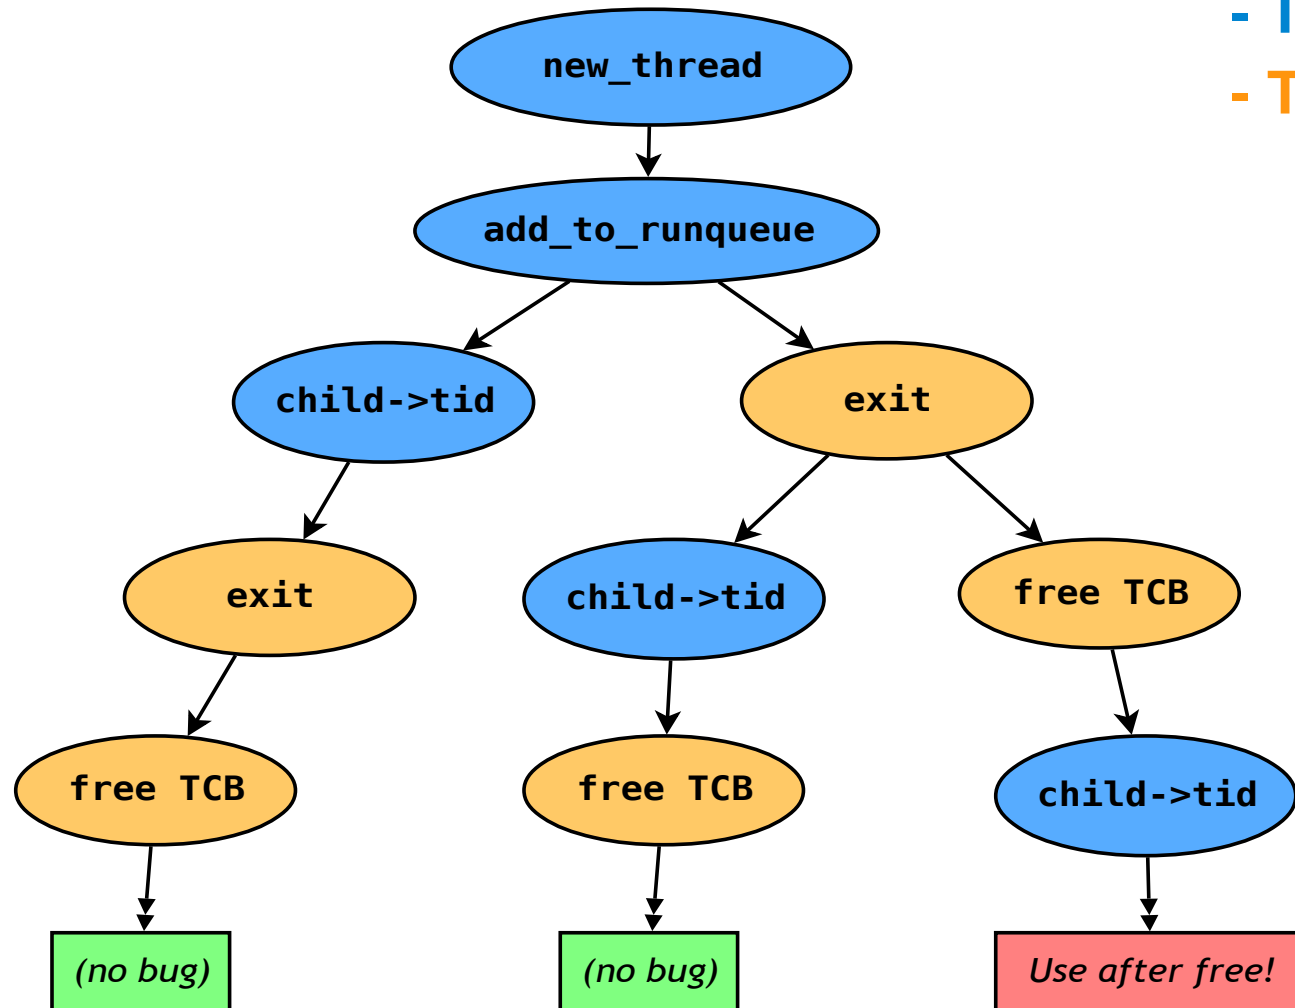
- Common approach to finding concurrency bugs
- Relies on randomly-occurring context switches
- User-friendly but unpredictable/unbounded

## Systematic testing [*Godefroid '97*]

- Thread scheduling controlled by test framework
- Each iteration, test a different thread interleaving
- Goal: Exhaustively search “all possible” interleavings
  - Wait... what?
  - In practice, interleavings are of coarse granularity.

# Example - State Space

- Thread 1
- Thread 2



# Preemption Points

---

The burning question of systematic testing:

“Which **preemption points** (PPs) are important?”

State space of interleavings is *parameterized* by PPs.

- Preemption points everywhere: completion infeasible
- Too few preemption points: won't find bugs

Prior work implementations hard-code a fixed set of PPs.


- dBug [Simsa '11]: pthread library calls
- Landslide [Blum '12]: kernel or user mutex lock/unlock

# Challenges

---

Parameters of systematic tests must be kept small.

- Test program length / number of threads
- Number of preemption points used



# Coping with Exponential Explosion

---

Some prior work focuses on improving *coverage* with...

- Reduction techniques (identifying equivalences)
  - *[Flanagan '05, Guo '11, Huang '15]*
- Determinizing runtimes *[Cui '13]*

Other prior systems give up on 100% completion...

- Random distribution of PPs *[Fonseca '14]*
- Heuristic exploration ordering (ICB) *[Musuvathi '08]*

# Coping with Exponential Explosion

---

## **Current systematic testing model not user-friendly.**

- Test framework: “I want to use these PPs, but can't predict how long until completion.”
- User: “I have 16 CPUs and 24h to test my program.”

Stress testing offers this... can we offer it too?

---

# Our Technique: Iterative Deepening

# Iterative Deepening

---

Goal: Run the best tests for a given CPU budget.

- User studies with 15-410 (operating systems)
  - Students worked best with an iterative process
  - “Start small, then add more preemption points”



# Iterative Deepening

---

Goal: Run the best tests for a given CPU budget.

- User studies with 15-410 (operating systems)
  - Students worked best with an iterative process
  - “Start small, then add more preemption points”

Named after analogous technique in chess AI.


- Chess search is DFS limited by max # of moves (ply).
- Chess AIs repeat DFS, increasing ply, until timeout.

# Iterative Deepening

---

“Minimal” state space == mandatory thread switches only


- `yield()`
- `cond_wait()`



# Iterative Deepening

Different PPs can produce state spaces of different sizes.

Testing them in parallel hedges our bets.




# Iterative Deepening

If time allows, we combine PPs to produce larger tests.

All PPs enabled = “maximal” state space

- Prior work tools explore this state space only.



# Choosing Preemption Points

---

How do we choose good preemption points?

- Start with “hardcoded” PPs as candidates
- Dynamic analysis during run can find new candidates:
  - Ad-hoc yield or busy loops
    - » Avoid cyclic state spaces
  - Data races [*Savage '97*]
    - » Including atomic operations enables testing locks, lock-free queues, etc.

Our main  
contribution

# Data-Race Preemption Points

---

**Data race:** 2 threads access the same memory, and:

- At least one access is a write
- Sets of locks held by each thread do not overlap
- No sync-enforced ordering (e.g. `cond_signal()`)

Let's preempt threads during these accesses!

- Preempting only on API boundaries can miss bugs
- 1-pass data-race analysis has false positive problem
  - Using data races as PPs automatically verifies or refutes each data race candidate.

# Landslide

---

**Landslide** [Blum '12]: simulator-based tester

- Originally built to test 15-410 Pebbles kernels
  - (a *landslide* tests the stability of *pebbles*)
- Wind River Simics: Full-system x86 simulator
  - Supports instruction/memory tracing, backtracking
  - No additional overhead for data race analysis

Landslide's features

- DPOR [Flanagan '05] prunes equivalences
- State space estimation guesses total time [Simsa '12]
- Bug detection: heap checking, deadlock, infinite loop...

# Landslide & Quicksand

---

**Quicksand:** wrapper program for Iterative Deepening

- Manages work queue of jobs with different PP configs
- Each job is a new state space for Landslide to explore
- Prioritizes which jobs are important / likely to finish
  - Based on nature of PPs (data races? mutexes?)
  - Based on estimated completion time

*Only required argument is user's CPU budget*



---

# Evaluation

# PLDI Evaluation Plan

---

Comparing to single-state-space (SSS) testing

- Finding bugs faster, while SSS times out
- Finding new bugs with data-race PPs

Comparing to 1-pass data-race analysis

- No details now, happy to explain at poster session

Testing OS projects from CMU, Berkeley, and U Chicago

- CMU: 79 “P2” thread libraries
- Berkeley & U Chicago: 79 “Pintos” kernels
- Under construction, feedback much appreciated!

# Test Suite

---

## 15-410 – Undergrad Operating Systems at CMU

- P2: pthread-like userspace thread library
  - Sync primitives: mutexes, cvars, sems, rwlocks
  - Thread API: `thr_create`, `thr_join`, `thr_exit`
- 6 test cases (50 LOC average):
  - *thr\_exit\_join, broadcast\_test, paraguay, rwlock\_test, paradise\_lost*
    - » Test of everything built on top of mutexes
  - *mutex\_test*
    - » Test of low-level lock implementation

# mutex test is special...

**Normally**, we give a lock's internal accesses a “free pass”.

- This enables productive data-race analysis elsewhere.

```
lock() { while (xchg(&lock->held, 1) == 1)  
        yield(lock->owner);  
        lock->owner = gettid();
```

```
        // in critical section...  
        // access protected data...
```

```
unlock() { lock->owner = -1;  
          lock->held = 0;
```

All accesses in  
this range are  
considered  
protected by  
the lock.

# mutex\_test is special...

---

In **mutex\_test**, we wish to verify the implicit assumption that other tests make that the locks are correct.

```
lock() { while (xchg(&lock->held, 1) == 1)
        yield(lock->owner);
        lock->owner = gettid();
```

```
    // in critical section...
    // access protected data...
```

```
unlock() { lock->owner = -1;
          lock->held = 0;
```

All accesses in this range are considered protected by the lock.

# Experimental Setup

---

79 P2s  $\times$  6 test cases  $\times$  10 CPU-hours  $\approx$  200 CPU-days

- 18 12-core (HT) 3.2 GHz Xeon machines, 12 GB RAM

Compare to a “control” experiment, also given 10 CPU-hours

- 1 state space, all “fixed” PPs enabled
  - mutex\_lock()
  - mutex\_unlock()
  - yield()
  - (cond\_wait(), etc, have mutex\_lock() PPs inside)

# Results

---

## Do we find the same bugs faster?

Control experiment with fixed PPs

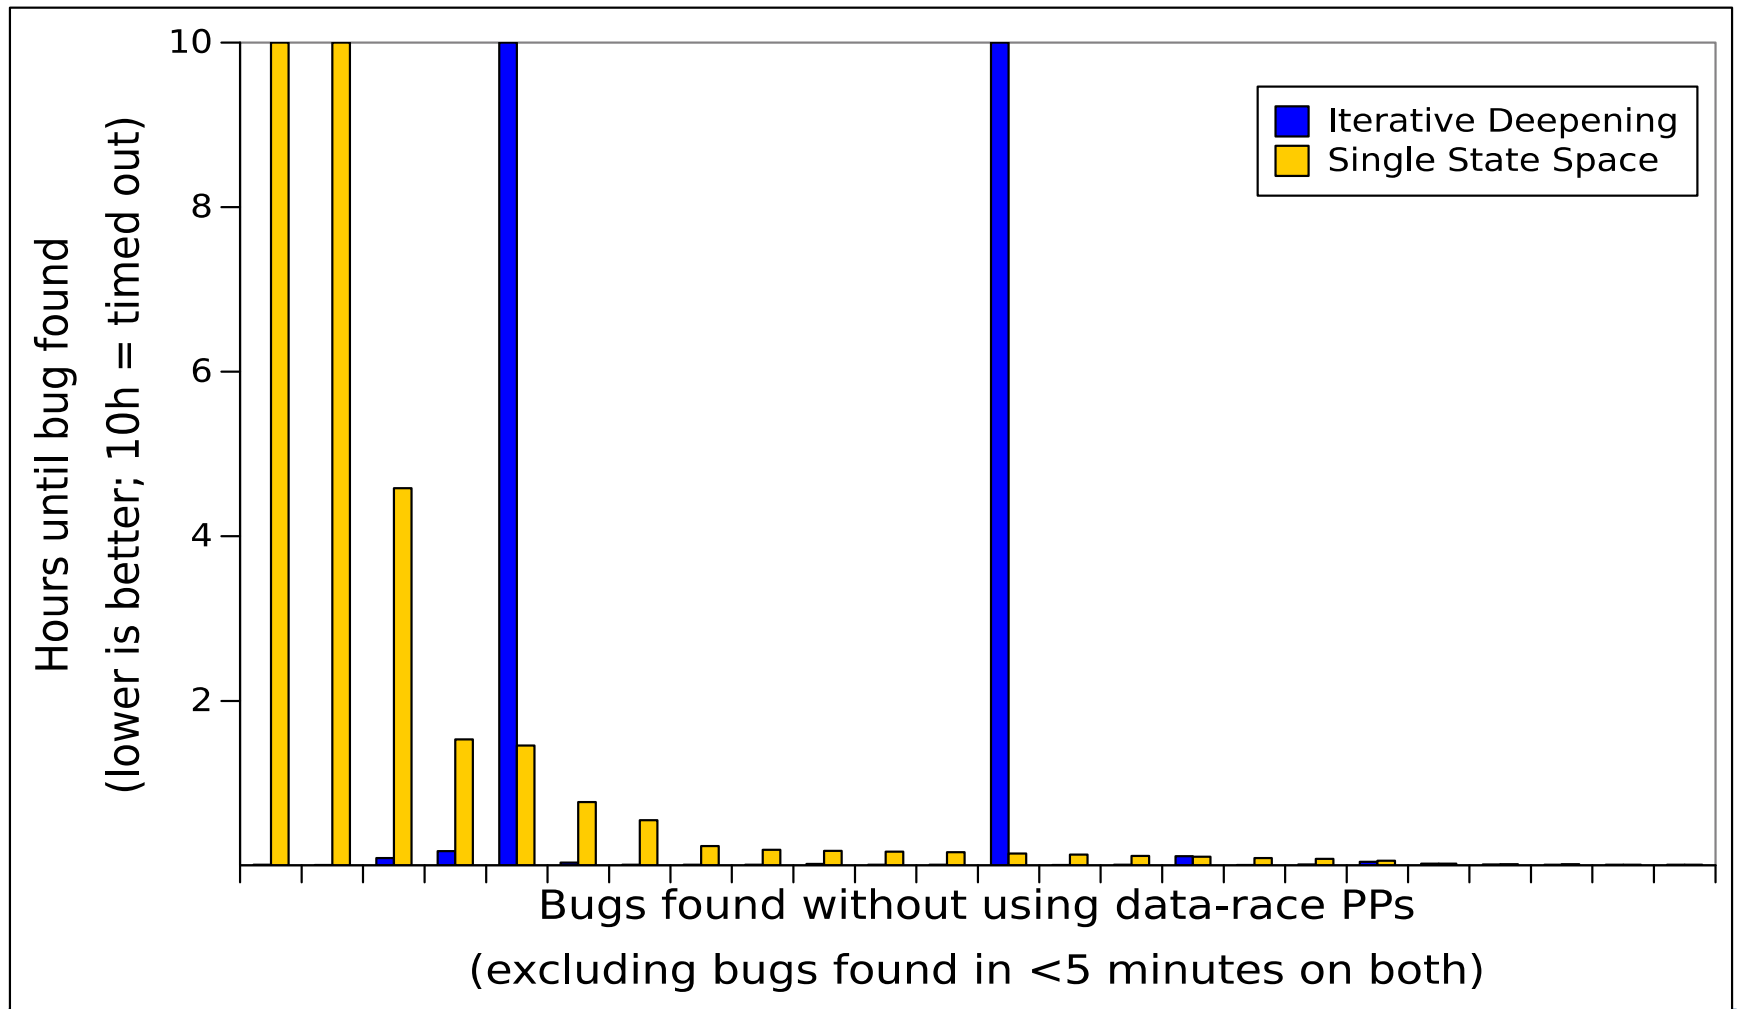
- 70 nondeterministic bugs
- 90 timeouts

Quicksand with iterative deepening

- Found 68 of those 70
  - Maximal state space's ETA was too high
- Found 2 more (no DR PPs needed) in control time-outs
  - Found much faster in subset state spaces

# Results

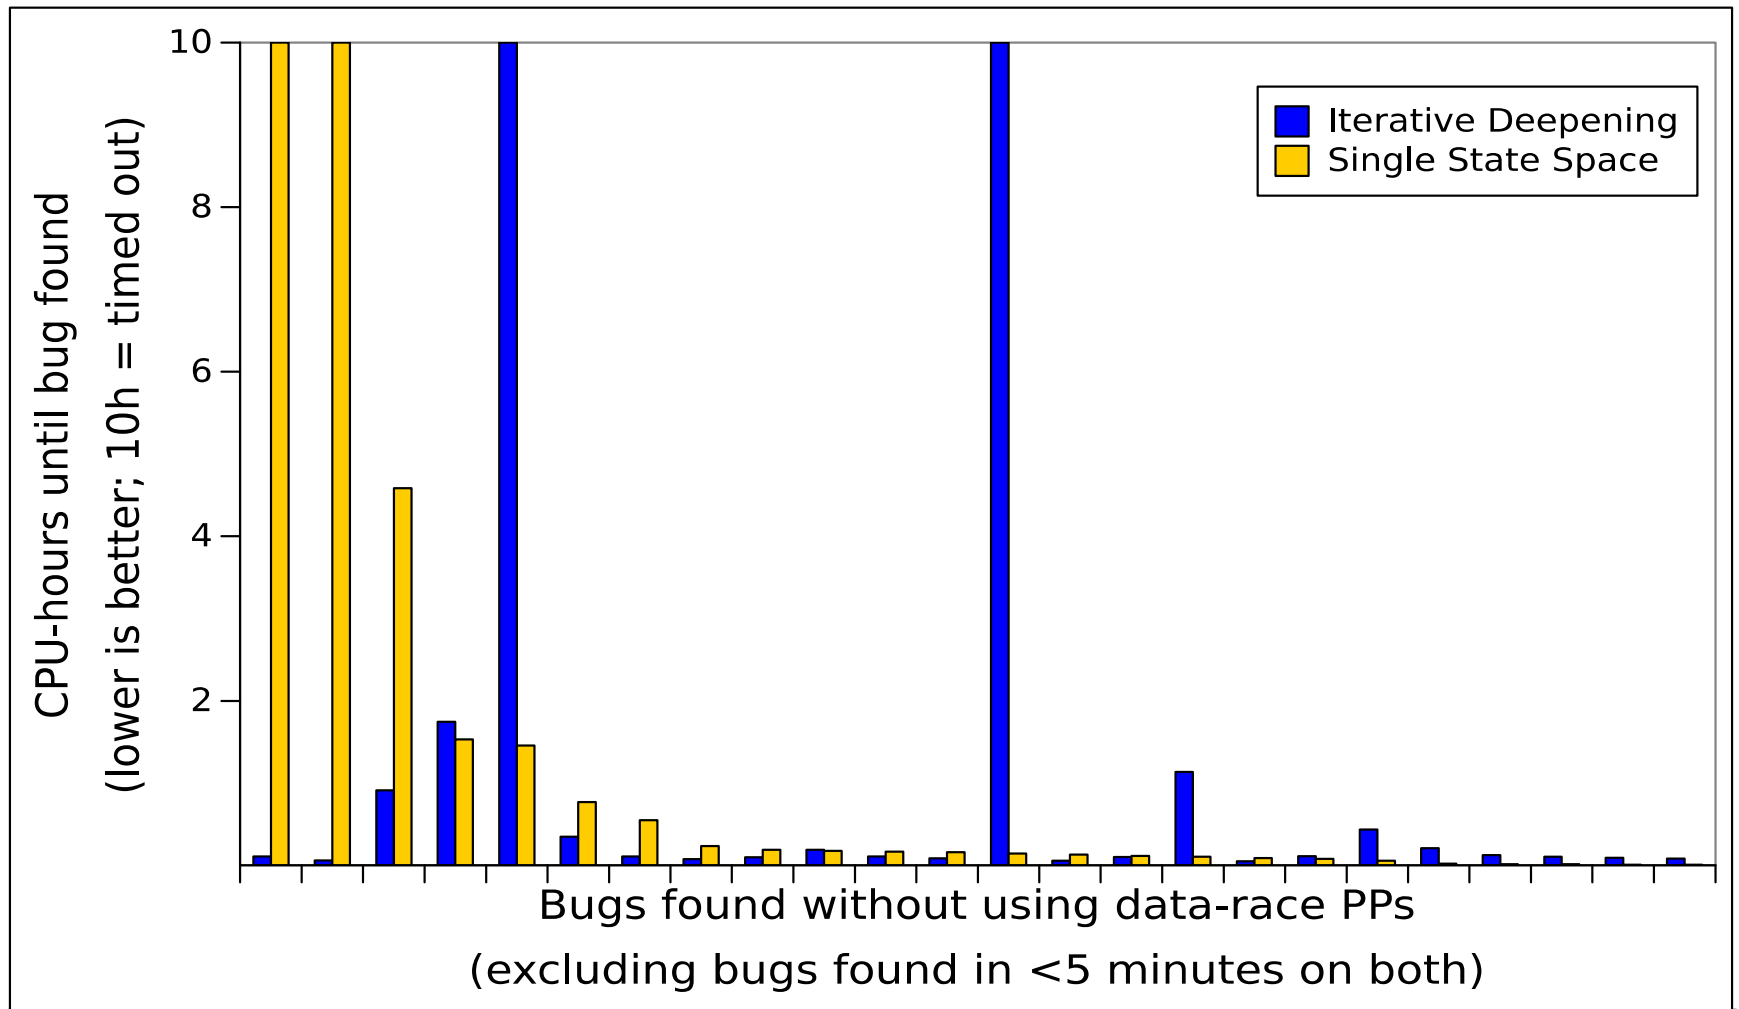
## Do we find the same bugs faster?





# Results

## Do we find the same bugs faster?



# Results

---

## Do we find new bugs with data-race PPs?

Quicksand with iterative deepening

- 107 nondeterministic bugs in total
- **33** required data-race PPs to expose!
  - 10 where control timed-out
  - 23 where control completed w/ no bug found
- `mutex_test`
  - 12 bugs (to control's 1)
  - All 12 others found with data-race PPs

---

# Conclusion

# Future Work

---

## Immediate future

- Ongoing user study with 15-410 students
  - Evaluate effectiveness for education/grading
- Finish PLDI evaluation plan, obviously
  - Have any experiment design ideas to share?

## More uncertain future

- Smarter policies for too-large state spaces (ICB)
- Incorporate parallel DPOR to saturate CPUs better
- Applying to production, not student, code

# Conclusion

---

## Systematic Testing

- Controls scheduling rather than rely on randomness
- More reliable bug-finding than stress testing
- State space depends on **preemption points** (PPs)
- Completion time is exponential and unpredictable

## Iterative deepening with Landslide & Quicksand

- Automatically searches for optimal set of PPs to use
- User need only supply CPU time budget
- Data race PPs uncover **50% more bugs**
- PLDI submission soon; want evaluation feedback

# References

---

## **[Godefroid '97]**

- Patrice Godefroid. VeriSoft: A Tool for the Automatic Analysis of Concurrent Reactive Software. CAV 1997.

## **[Savage '97]**

- Stefan Savage et al. Eraser: A Dynamic Data Race Detector for Multithreaded Programs. ACM TOCS 1997.

## **[Flanagan '05]**

- Cormac Flanagan and Patrice Godefroid. Dynamic partial-order reduction for model checking software. POPL 2005.

## **[Musuvathi '08]**

- Madanlal Musuvathi et al. Finding and Reproducing Heisenbugs in Concurrent Programs. (a.k.a. the CHES paper.) OSDI 2008.

## **[Simsa '11]**

- Jirí Simsa, Randy Bryant, Garth A. Gibson: dBug: Systematic Testing of Unmodified Distributed and Multi-threaded Systems. SPIN 2011.

# References

---

## [Guo '11]

- Huayang Guo et al. Practical Software Model Checking via Dynamic Interface Reduction. SOSP 2011.

## [Blum '12]

- Ben Blum. Landslide: Systematic Dynamic Race Detection in Kernel Space. CMU-CS-12-118. May 2012.

## [Simsa '12]

- Jiri Simsa. Runtime Estimation and Resource Allocation for Concurrency Testing. CMU-PDL-12-113. December 2012.

## [Cui '13]

- Heming Cui et al. Parrot: A Practical Runtime for Deterministic, Stable, and Reliable Threads. SOSP 2013.

## [Fonseca '14]

- Pedro Fonseca et al. SKI: Exposing Kernel Concurrency Bugs through Systematic Schedule Exploration. OSDI 2014.

# Related Work

---

## Systematic testing

- MaceMC (NSDI '07) – liveness, random walking
- CHES (PLDI '07) – iterative context bounding (ICB)
- MoDist (NSDI '09) – network/disk model checking
- dBug (SSV '10/SPIN '11) – dynamic partial order reduction
- SimTester (VEE '12) – interrupt injection, drivers
- Parrot (SOSP '13) – dBug + deterministic userspace scheduler
- SKI (OSDI '14) – randomly-chosen PPs for testing Linux
- SAMC (OSDI '14) – annotated sync primitives, better reduction

## Data race detection

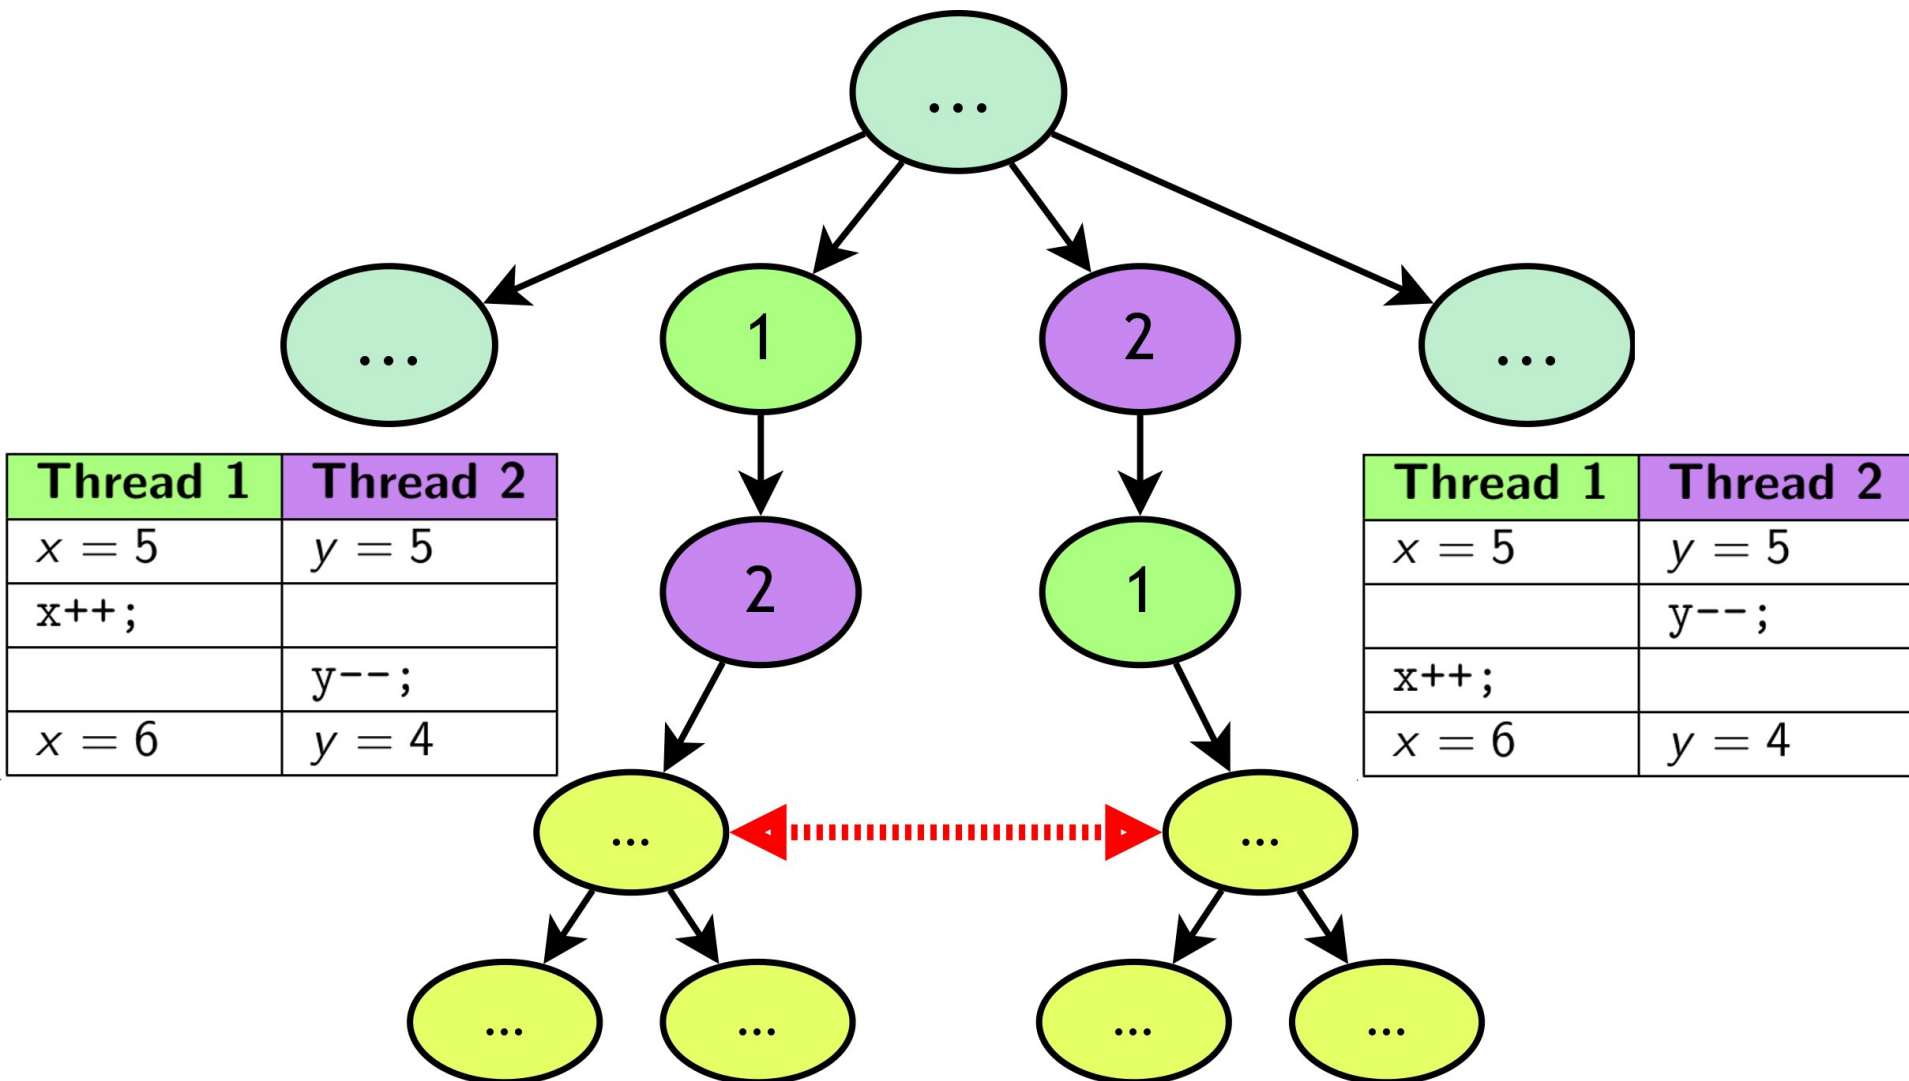
- Eraser (TOCS '97) – lock-set tracking, annotations
- DataCollider (OSDI '10) – random sampling, kernel
- RacePro (SOSP '11) – inter-process races
- Portend (ASPLOS '12) – alternate interleavings, symbolic execution



---

# Bonus Slides

# Dynamic Partial Order Reduction



# Results

---

Among 474 total tests (P2+unit test pairs)...

- 23 *deterministic* bugs (e.g. use-after-free); fixed by hand

Control experiment with fixed PPs

- 70 nondeterministic bugs
- 90 timeouts

Quicksand with iterative deepening

- 107 nondeterministic bugs
  - 37 requiring data-race PPs to expose
  - mutex\_test: 13 bugs, 12 requiring data-race PPs
- 2 bugs found without data-race PPs (control timeout)
- 2 bugs missed that control found (*explain why*)