

---

# **Landslide:** Systematic Testing for Kernel Space Race Detection

**Ben Blum**

**Jiří Šimša, Garth Gibson**

Parallel Data Laboratory  
Carnegie Mellon University

# Outline

---

## **Motivation: Concurrency debugging**

- Systematic testing versus stress testing
- State space explosion
- Challenges of kernel-space

## **Tool: Landslide**

- Design and interface
- Addressing challenges

## **Evaluation: Finding Races**

- Student user study
- Case study

## **Future Work and Conclusion**

# Motivation – Example

---

```
int thread_fork()  
{  
    thread_t *child = spawn_new_thread();  
    add_to_runqueue(child);  
    return child->tid;  
}
```

# Motivation – Example

---

```
int thread_fork()  
{  
    thread_t *child = spawn_new_thread();  
    add_to_runqueue(child);  
    return child->tid; ← “child” gets freed!  
}
```

- On exit, child's state is freed
- Forking thread does use-after-free
- Might return garbage instead of thread ID

# Motivation – Testing Techniques

---

## **Stress Testing:** Common testing approach

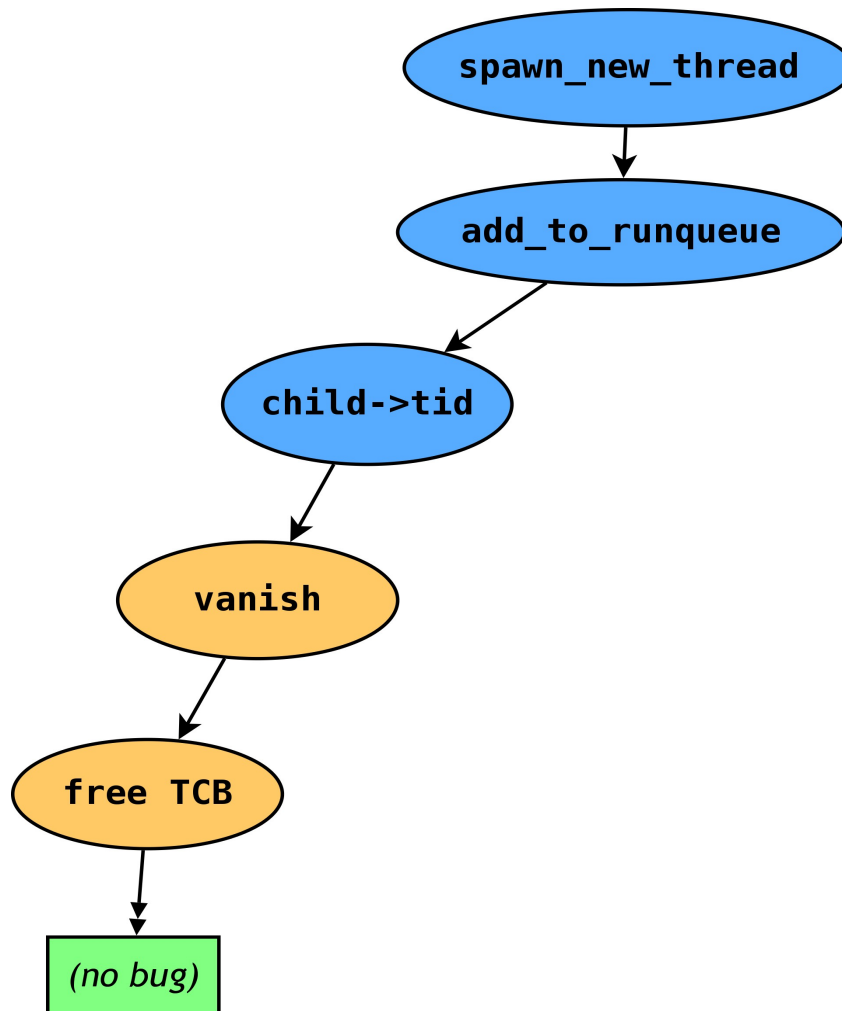
- Trying to exercise as many interleavings as possible
- Exposes race conditions at random
  - “If a preemption occurs at *just the right time*...”
- Cryptic panic messages or machine reboots

## **Systematic Testing** [*Godefroid '97*]

- Make educated guesses about when to preempt
- Run *every single* interleaving
- Provide better debugging information, reproducibility

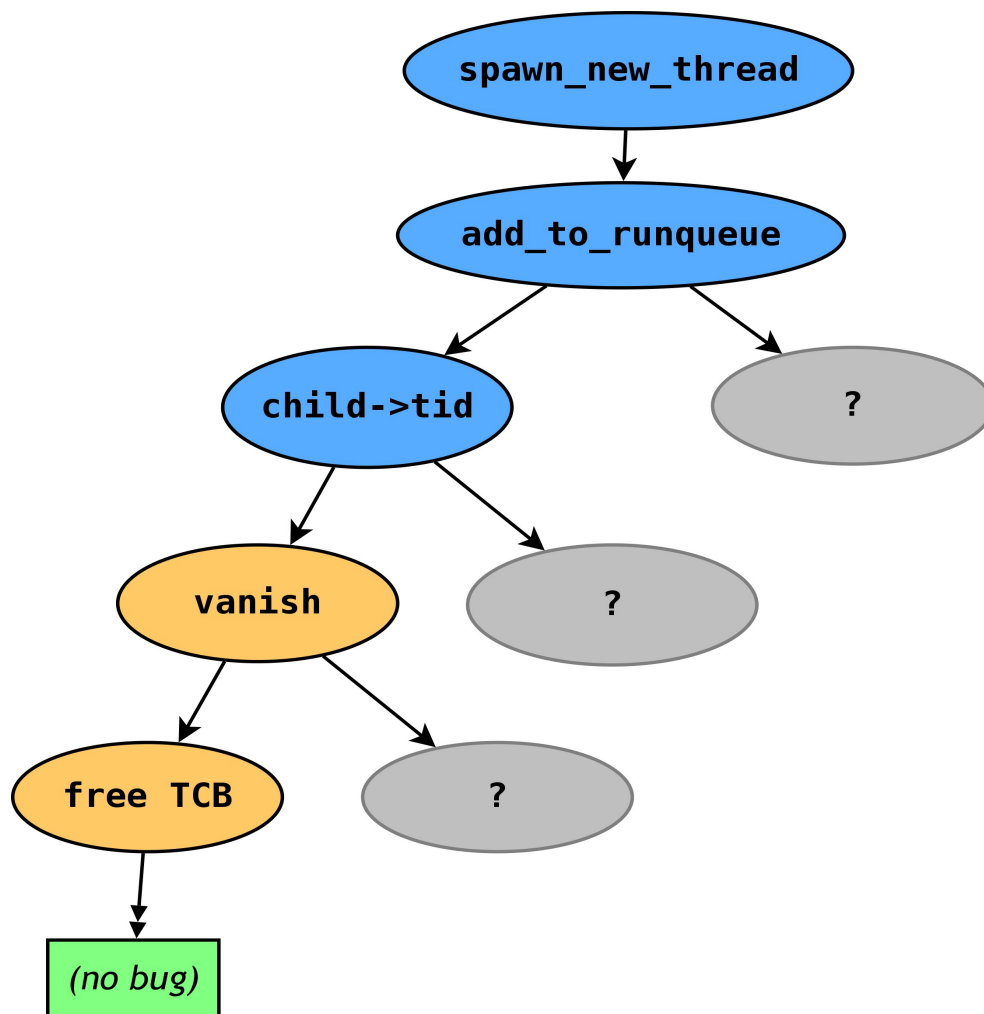
# The Execution Tree

---

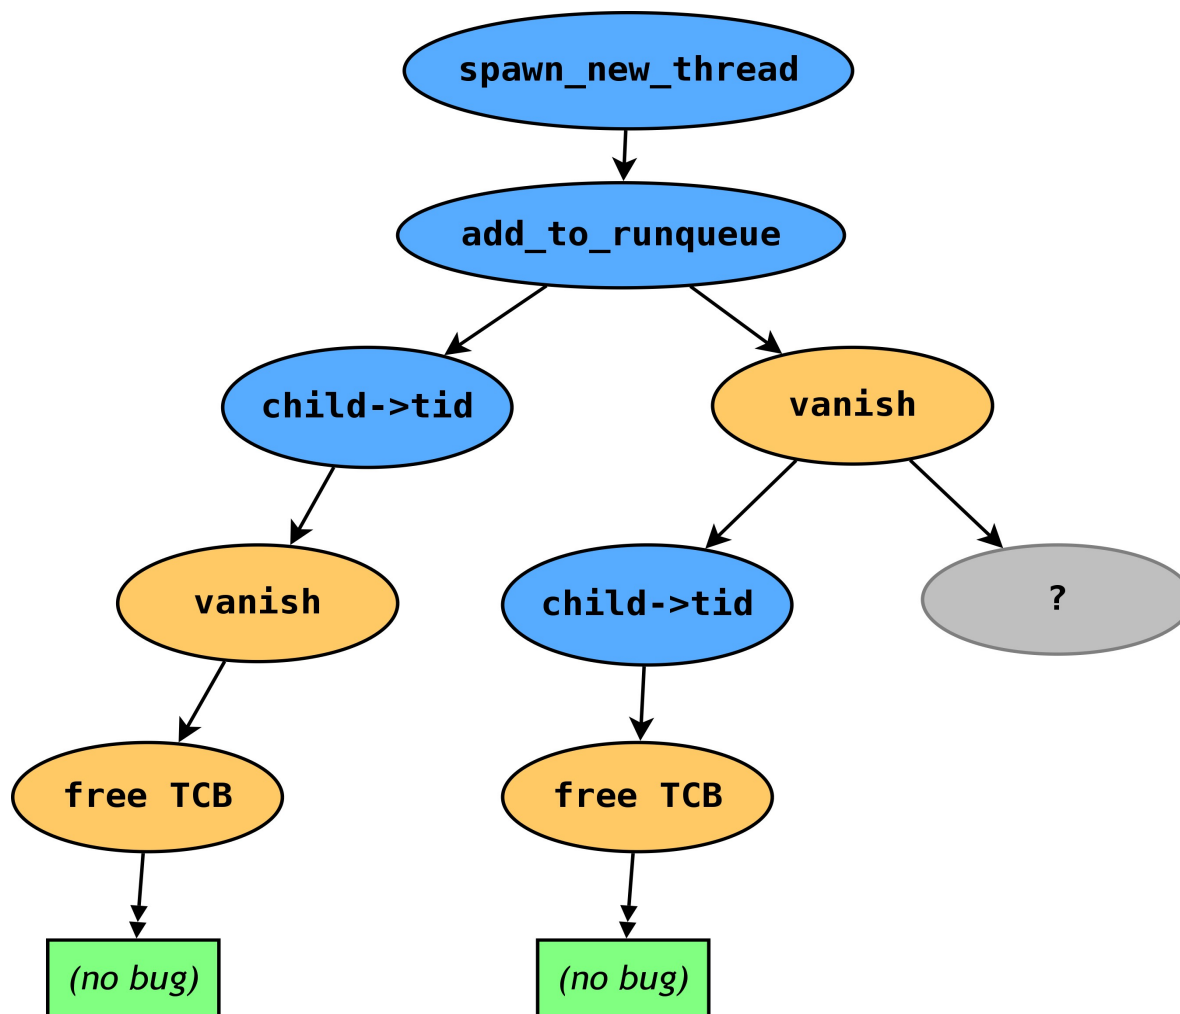


# The Execution Tree

---

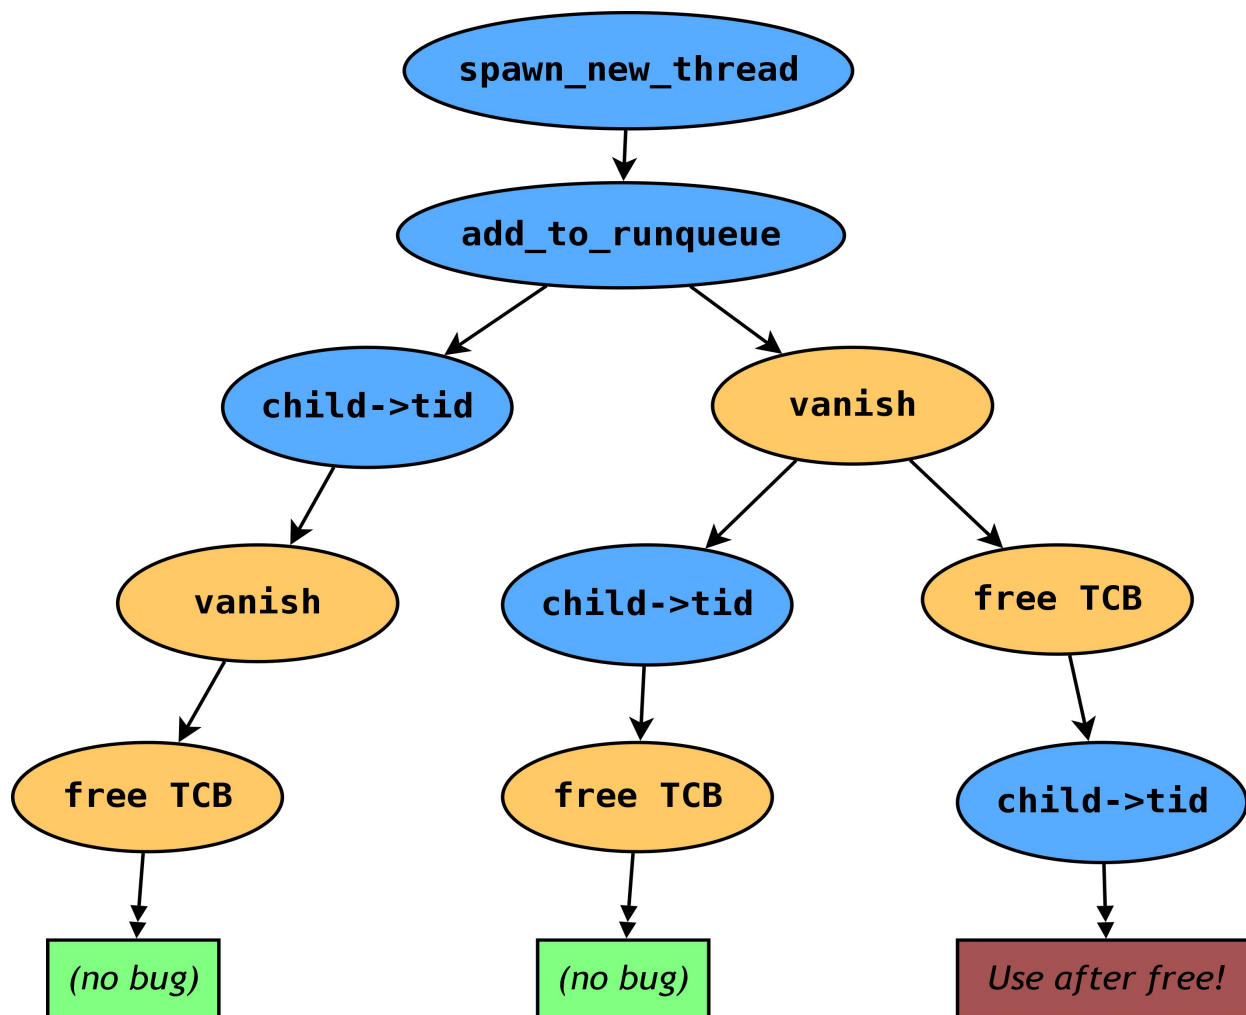


# The Execution Tree





# The Execution Tree



# Decision Points

---

**Decision points:** where being preempted causes different behaviour.

**What does “all possible interleavings” mean?**

- One extreme: Decide at every instruction
  - Impossibly large test runtime
- Other extreme: No decision points
  - Makes “no race found” a weak claim
- Sweet spot: Intuit where preemptions “might matter”
  - Joint effort between programmer and tool

# Decision Points

---

**In userspace, pthread library calls. *[Simsa '11]***

**In kernels, we need to tell the story from scratch.**

# Challenge 1: Causes of Concurrency

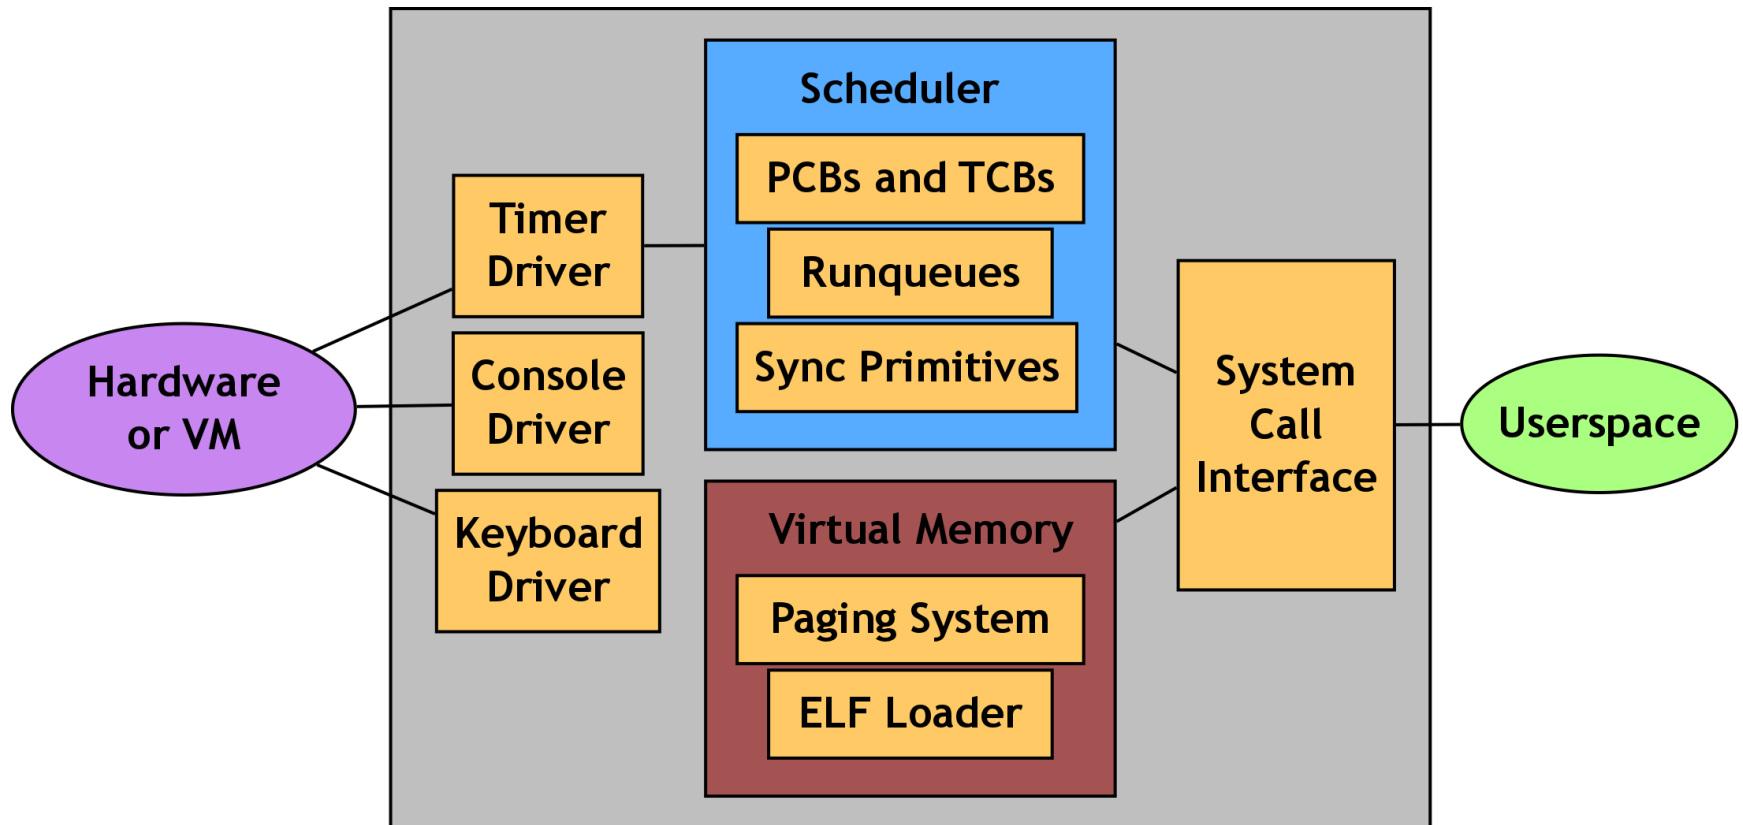
---

## **Kernels contain their own concurrency implementation.**

- Context switching
  - Nondeterministic timer-driven thread scheduling
- Runqueue tracking
  - Which threads are runnable?
- Thread lifecycle tracking
  - When are threads created/destroyed?

# Challenge 2: The Kernel as One Program

**“Everything interleaves with everything else”?**



# Contribution

---

**Systematic testing is a *tool* for testing concurrent systems.**

- How can it be applied in kernel-space?
- What simple things can the user provide to help?

# Outline

---

## *Motivation: Concurrency debugging*

- *Systematic testing versus stress testing*
- *State space explosion*
- *Challenges of kernel-space*

## **Tool: Landslide**

- Design and interface
- Addressing challenges

## *Evaluation: Finding Races*

- *Student user study*
- *Case study*

## *Future Work and Conclusion*

# Execution Environment

---

**Simics:** a full-system x86 simulator

- Landslide runs as a Simics module.
  - Can see kernel instructions, memory accesses

## Controlling system nondeterminism

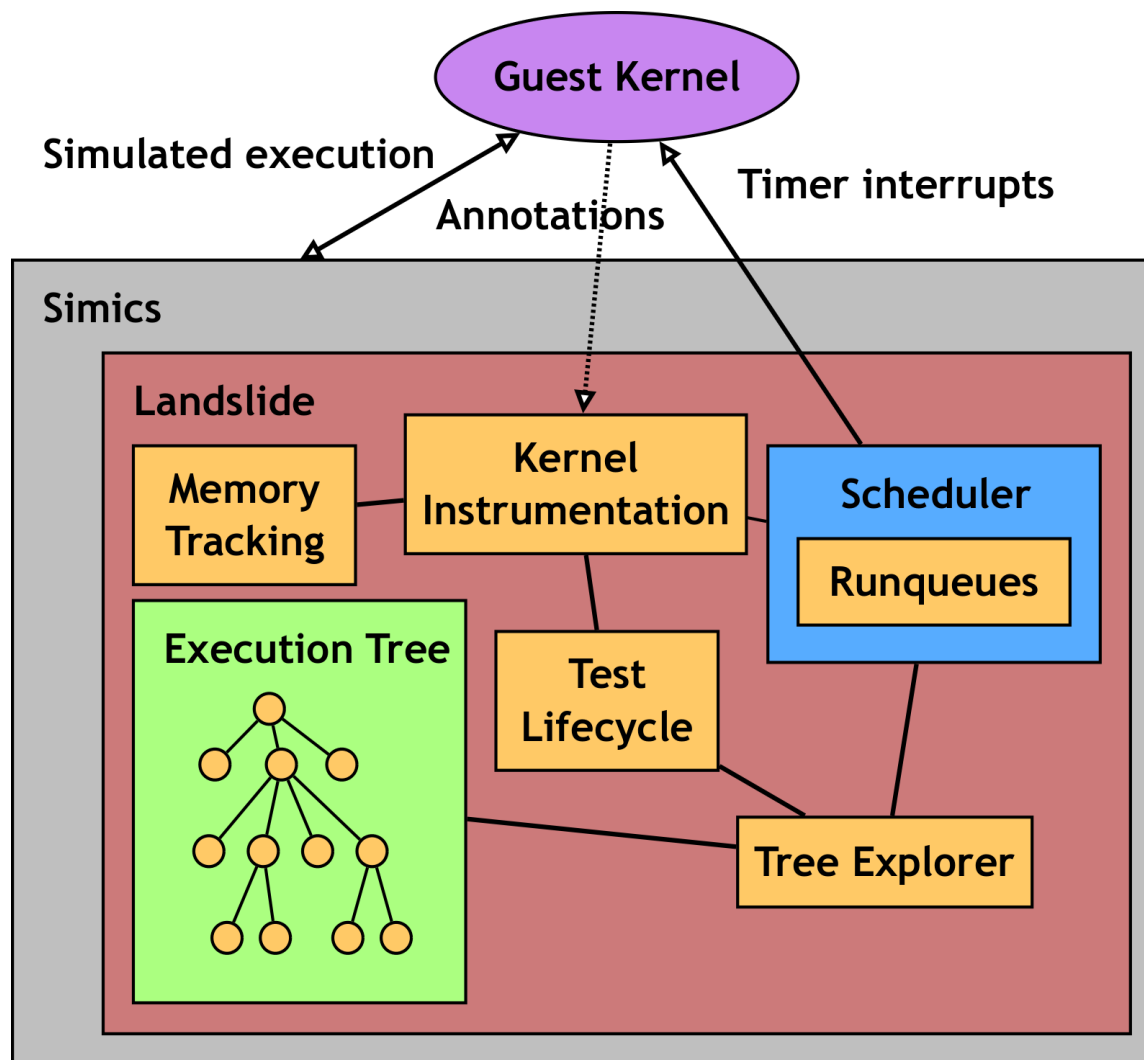
- Injecting timer interrupts triggers thread scheduling
- Future work may use device interrupts (disk, network)

## Backtracking

- After each test execution, can rewind system state using Simics “bookmarks”



# Anatomy of Landslide



# Interface

---

```
int thread_fork()
{
    thread_t *child = spawn_new_thread();
    tell_landslide_forking();
    add_to_runqueue(child);
    tell_landslide_decide(); /* Interrupt me here! */
    return child->tid;
}
```

## User must tell Landslide when certain events happen:

- When do threads become runnable / descheduled?
- When does the scheduler switch threads?

## Can also configure optional additional decision points

# Decision Trace

---

**USE AFTER FREE - read from 0x0015a8f0 at PC 0x00104209**  
[0x15a8f0 | 4136] was allocated by TID3 at (...)  
and freed by TID4 at (...)

\*\*\*\* A bug was found! \*\*\*\*

\*\*\*\* Decision trace follows. \*\*\*\*

**1: 1347079 instructions, old TID 3, new TID 4**

TID3 at 0x00105a10 in **context\_switch**,  
0x001041f4 in **thread\_fork**,  
0x0010362b in **thread\_fork\_wrapper**

**2: 1350725 instructions, old TID 4, new TID 3**

TID4 at 0x00105a10 in **context\_switch**,  
0x00104681 in **yield**,  
0x00104570 in **exit**,  
0x00103708 in **exit\_wrapper**

Stack: TID3 at 0x00104209 in **thread\_fork**,  
0x0010362b in **thread\_fork\_wrapper**

**Total decision points 24, total backtracks 5**

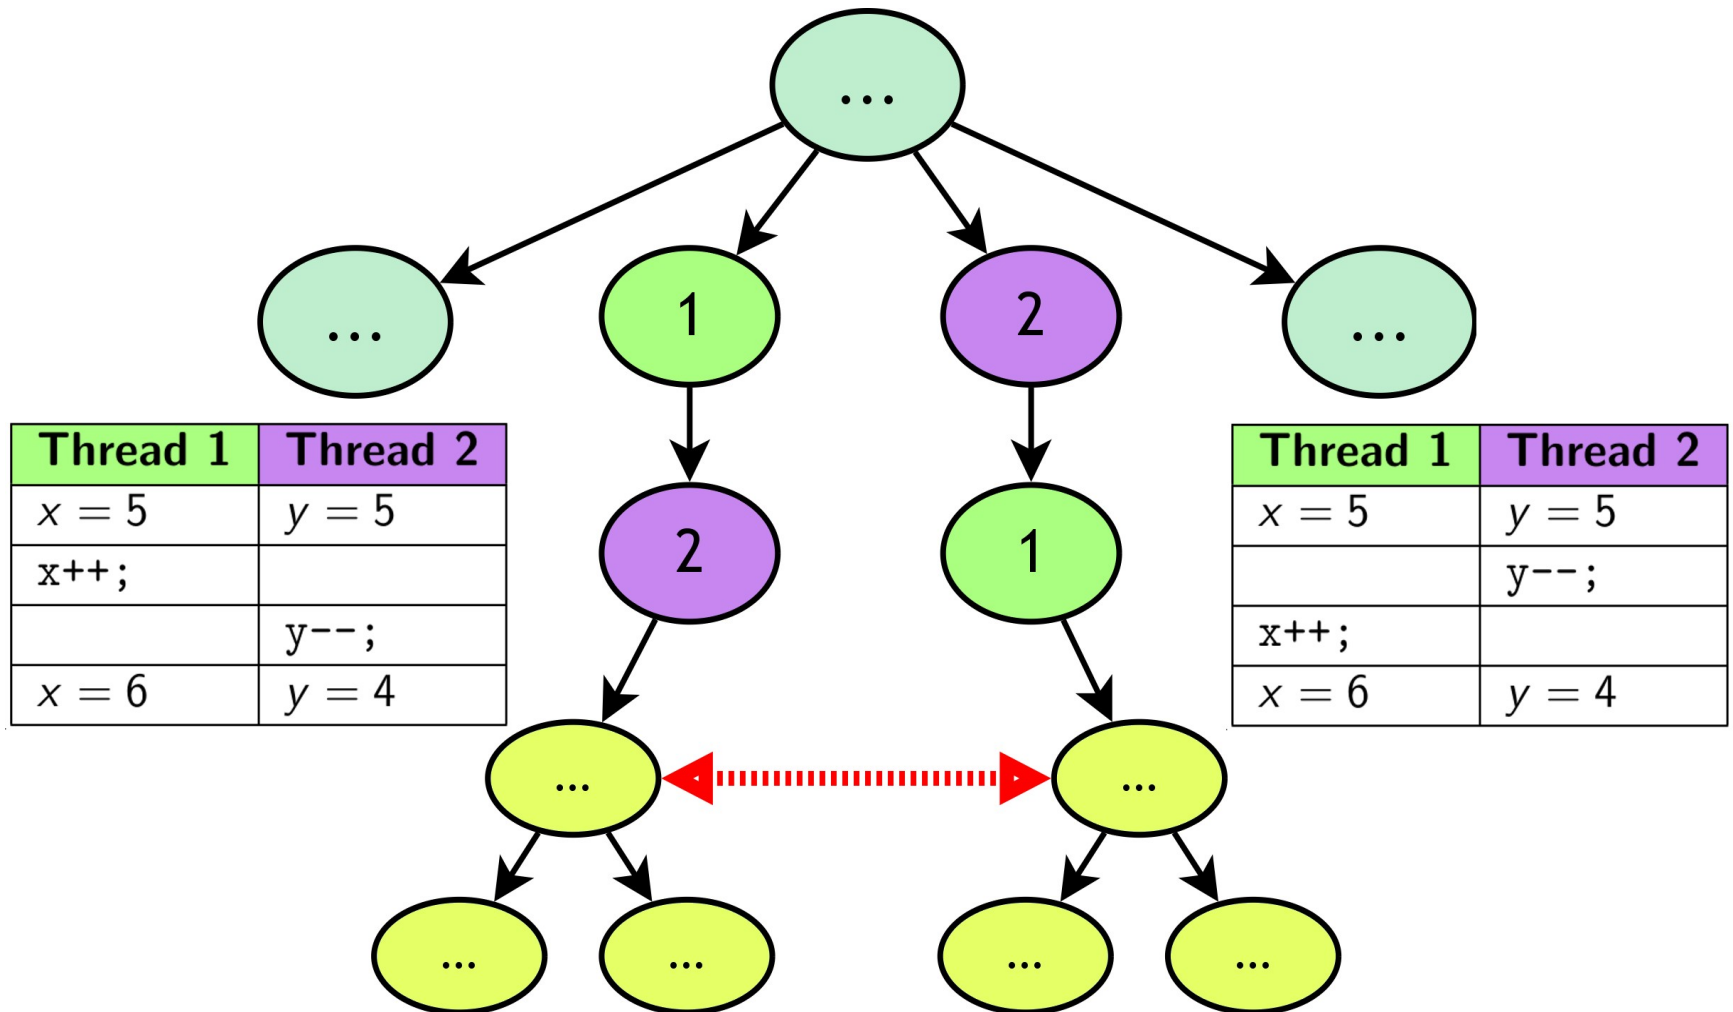
# State Space Reduction

---

## State spaces grow exponentially.


- Fortunately, some sequences cause identical states.
- Dynamic Partial Order Reduction [*Flanagan '05*]
  - Can even be parallelized [*Simsa '12*]
  - Requires “memory independence relation” between transitions

# State Space Explosion



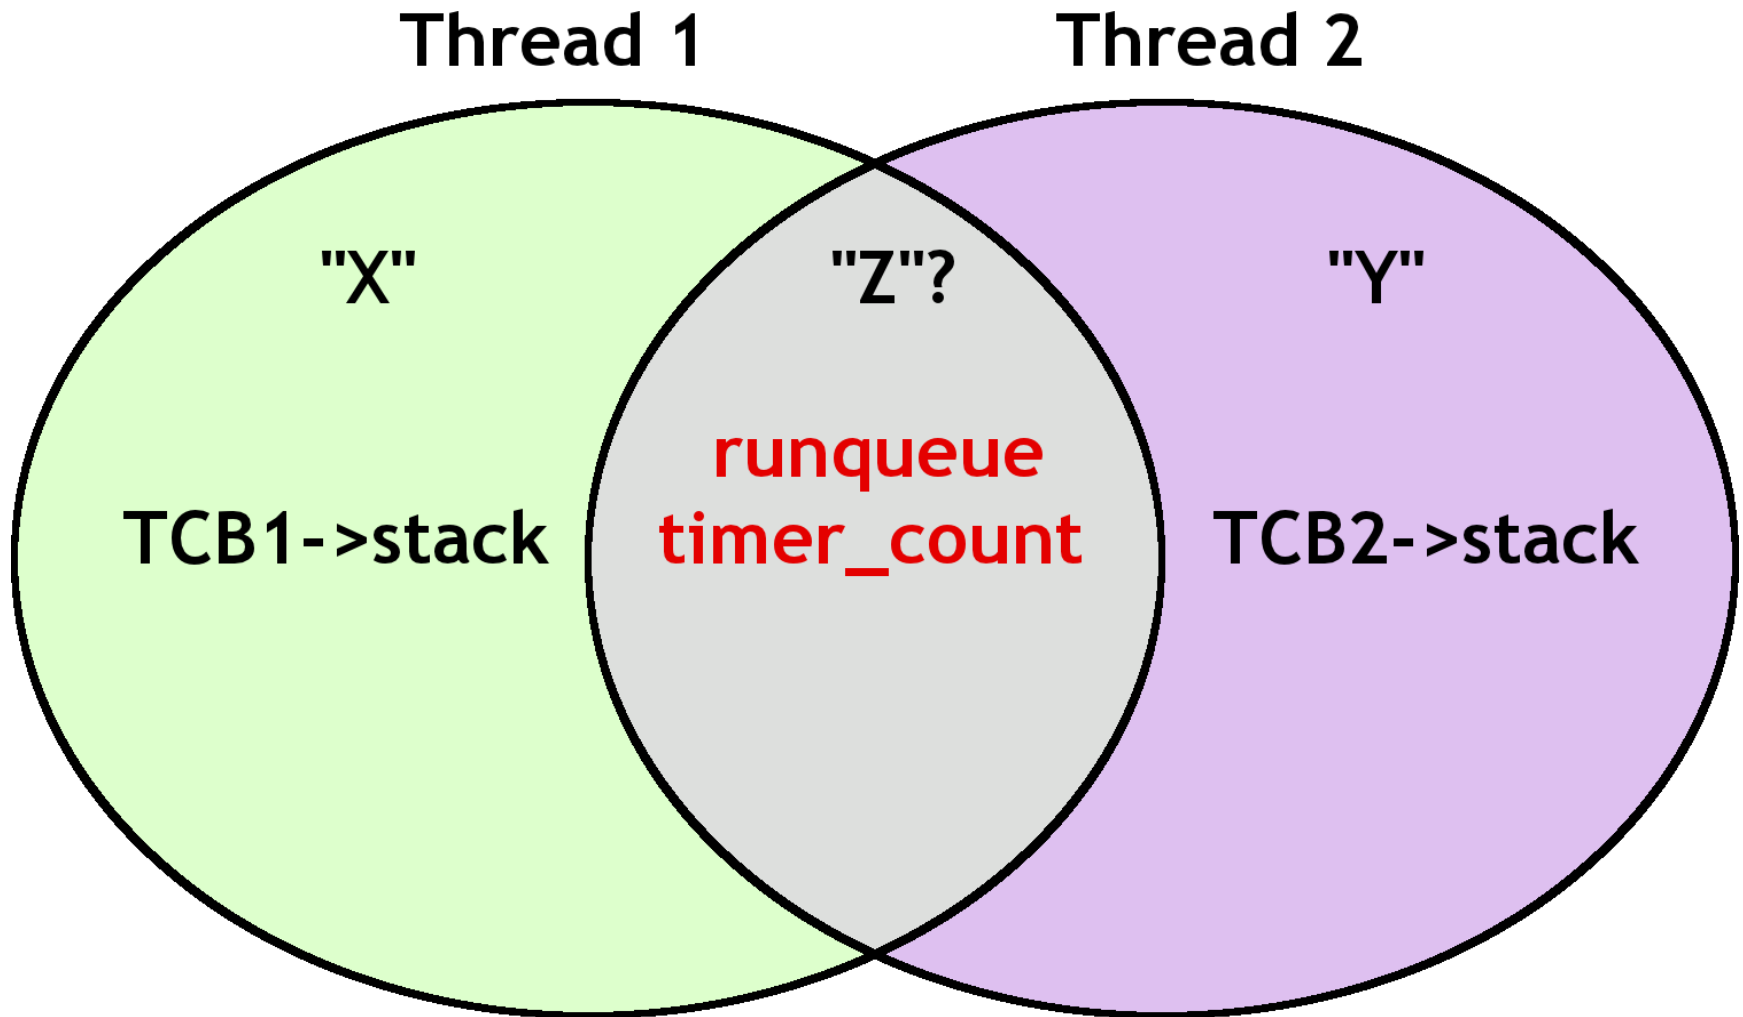
# Shared Memory Conflicts

---



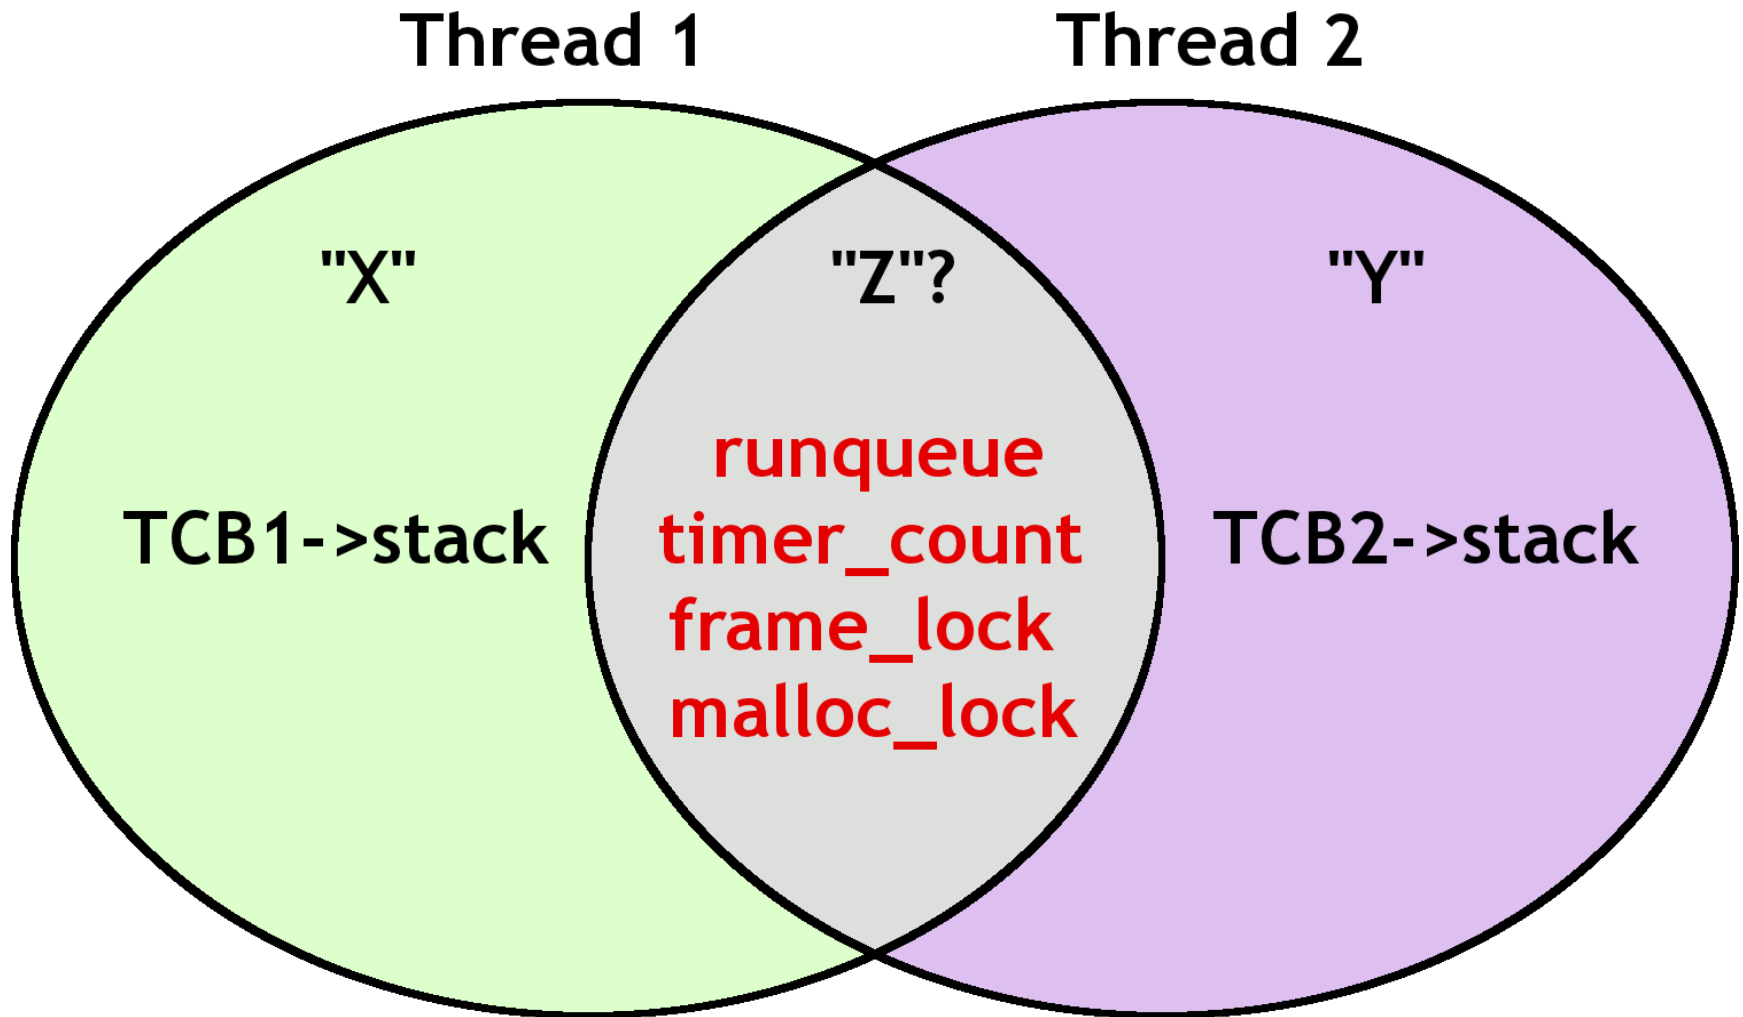
# Shared Memory Conflicts

---



# Shared Memory Conflicts

---





# Shared Memory Conflicts

---

**Solution:** Only consider “relevant” memory conflicts.

- Ignore scheduler, global objects assumed to be correct.
  - Sacrifice ability to test these to efficiently test everything else.
- User must specify what to ignore.

# Focusing the Search Space

---

## Recommended decision points:

`mutex_lock()`, `mutex_unlock()`

**Problem:** Kernel uses mutexes *everywhere*.

- Good: `exit()` calls `mutex_lock()`
- Bad: `destroy_address_space()` calls it too

**Solution:** User specifies which modules to pay attention to.

- `within_function` `exit`
- `without_function` `destroy_address_space`

# Identifying Bugs

---

**How do we know we've found a bug?**

## **Definite bugs**

- Kernel panics / assertion failures
- Use-after-free
- Deadlock

## **Probable bugs**

- Infinite loop
  - Use structure of execution tree to judge progress.

# Identifying Bugs


**How do we know we've found a bug?**

## Definite bugs

- Kernel panics / assertion failures
- Use-after-free
- Deadlock

## Probable bugs

- Infinite loop
  - Use structure of execution tree to judge progress.



# Outline

---

## *Motivation: Concurrency debugging*

- *Systematic testing versus stress testing*
- *State space explosion*
- *Challenges of kernel-space*

## *Tool: Landslide*

- *Design and interface*
- *Addressing challenges*

## **Evaluation: Finding Races**

- Student user study
- Case study

## *Future Work and Conclusion*

# Working with Students

---

## **15-410: Operating System Design and Implementation**

- Students implement a small kernel in 6 weeks
- “Pebbles”, a UNIX-like system call specification
- On average 4000 lines of code

## **Solicited students to use Landslide on their kernels.**

- Can Landslide find bugs “in the wild”?
- How much time does manual instrumentation take?

## **Five groups participated; four got it to work.**

- Average instrumentation time 100 minutes (55-158)
- All four groups found bugs; two were nondeterministic

# Studying Bugs In-Depth


---

## **Tested two Pebbles kernels: my own, and one I graded**

- Confirmed several of the most subtle races I had found in the kernel I graded
  - Parent and child exit simultaneously
- Found a previously unknown bug in my kernel
  - Condition variable wait queue corruption
  - Missed by TA's manual inspection, stress tests

# Studying Bugs In-Depth

**Exploration ordering: Search branches with more or fewer preemptions first?**






# Studying Bugs In-Depth


## Exploration ordering

- “Backwards” exploration found bugs sooner
- “Forwards” exploration found shorter decision traces

Exploration Time



Decision Trace Length



# Outline

---

## *Motivation: Concurrency debugging*

- *Systematic testing versus stress testing*
- *State space explosion*
- *Challenges of kernel-space*

## *Tool: Landslide*

- *Design and interface*
- *Addressing challenges*

## *Evaluation: Finding Races*

- *Student user study*
- *Case study*

## **Future Work and Conclusion**

# Future Work

---

## Performance

- Virtualization – less control, compared to Simics

## Generalizing the concurrency model

- Device driver input as new source of nondeterminism
- Multicore execution
- Better model more complex kernels
  - e.g. educational (Pintos) or production (Linux)

## Emphasis on “steering” test by changing parameters

- Start with coarse granularity, iteratively refine
- Perhaps we can automate this steering?

# Conclusion

---

## **Systematic testing**

- Make educated guesses about when to preempt.
- Find many types of races; provide debugging info.

## **Systematic testing in kernel-space**

- Use internal kernel abstractions to understand concurrency behaviour.
- Relying on user's knowledge makes testing easier.

## **Landslide**

- A first step towards sophisticated kernel debugging techniques.

# Related Work

---

## Systematic testing

- MaceMC (NSDI '07) – liveness, random walking
- CHES (PLDI '07) – iterative context bounding
- MoDist (NSDI '09) – network/disk model checking
- dBug (SSV '10) – dynamic partial order reduction
- SimTester (VEE '12) – interrupt injection, drivers

## Data race detection

- Eraser (TOCS '97) – lock-set tracking, annotations
- DataCollider (OSDI '10) – random sampling, kernel
- RacePro (SOSP '11) – inter-process races

# References

---

## **[Godefroid '97]**

- Patrice Godefroid. VeriSoft: A Tool for the Automatic Analysis of Concurrent Reactive Software. CAV 1997.

## **[Flanagan '05]**

- Cormac Flanagan and Patrice Godefroid. Dynamic partial-order reduction for model checking software. POPL 2005.

## **[Simsa '11]**

- Jiri Simsa, Randy Bryant, Garth A. Gibson: dBug: Systematic Testing of Unmodified Distributed and Multi-threaded Systems. SPIN 2011.

## **[Simsa '12]**

- Jiri Simsa. Scalable Dynamic Partial Order Reduction. Talk at PDL Retreat 2012.